

assignment_1_handout

November 8, 2024

1 Assignment 1 - Linear Regression

Submitted by: - Jonas Henker, 7054995 - Leo Forster, 7055800

In this assignment you will be coding for a Linear Regression task hands-on. (10 Points)

The notebook uses some popular libraries. If your environment is missing any of these libraries, you can install them using the following pip commands:

“bash !pip install matplotlib seaborn scikit-learn

```
[1]: import math

from sklearn.datasets import fetch_california_housing
import pandas as pd
from pandas.plotting import scatter_matrix
from scipy import stats
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

```
[2]: #make sizes bigger for readability
import matplotlib.pyplot as plt
plt.rcParams.update({'font.size': 17})
plt.rcParams["figure.figsize"] = (12,12)
```

1.1 Load and Explore Data

```
[3]: # Load the California Housing dataset
housing = fetch_california_housing()
# Convert the dataset into a DataFrame
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['MedHouseVal'] = housing.target # Add the target (median house value)
```

Number of Instances:

20640

Number of Attributes:

8 numeric, predictive attributes and the target

Attribute Information:

MedInc median income in block group

HouseAge median house age in block group

AveRooms average number of rooms per household

AveBedrms average number of bedrooms per household

Population block group population

AveOccup average number of household members

Latitude block group latitude

Longitude block group longitude

[4]: display(df)

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	
...	
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807	39.49	
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635	39.43	
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981	39.37	

	Longitude	MedHouseVal
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422
...
20635	-121.09	0.781
20636	-121.21	0.771
20637	-121.22	0.923
20638	-121.32	0.847
20639	-121.24	0.894

[20640 rows x 9 columns]

```
[5]: #Explore data for missingness
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MedInc           20640 non-null  float64
1   HouseAge         20640 non-null  float64
2   AveRooms         20640 non-null  float64
3   AveBedrms        20640 non-null  float64
4   Population       20640 non-null  float64
5   AveOccup         20640 non-null  float64
6   Latitude         20640 non-null  float64
7   Longitude        20640 non-null  float64
8   MedHouseVal      20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
None
```

```
[6]: #Check statistics of the data
print(df.describe())
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	\
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	
std	1.899822	12.585558	2.474173	0.473911	1132.462122	
min	0.499900	1.000000	0.846154	0.333333	3.000000	
25%	2.563400	18.000000	4.440716	1.006079	787.000000	
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	
max	15.000100	52.000000	141.909091	34.066667	35682.000000	

	AveOccup	Latitude	Longitude	MedHouseVal
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704	2.068558
std	10.386050	2.135952	2.003532	1.153956
min	0.692308	32.540000	-124.350000	0.149990
25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010

```
[7]: # Display the first few rows
print(df.head())
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	

1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85

	Longitude	MedHouseVal
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422


```
[8]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Select multiple features for the correlation check
X_all = df[['MedInc', 'AveRooms', 'AveOccup', 'HouseAge', 'Latitude',
            ↪ 'Longitude']]

# Calculate correlation matrix
corr_matrix = X_all.corr()

# Visualize the correlation matrix
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Matrix of Features')
plt.show()

# Note that correlation between Latitude and Longitude is coming from ↪
↪ geographical location of California
```



```
[9]: #display scatter_matrix also
fig = plt.figure()
scatter_matrix(df,figsize =(25,25),alpha=0.9,diagonal="kde",marker="o")
```

```
[9]: array([[<Axes: xlabel='MedInc', ylabel='MedInc'>,
<Axes: xlabel='HouseAge', ylabel='MedInc'>,
<Axes: xlabel='AveRooms', ylabel='MedInc'>,
<Axes: xlabel='AveBedrms', ylabel='MedInc'>,
<Axes: xlabel='Population', ylabel='MedInc'>,
<Axes: xlabel='AveOccup', ylabel='MedInc'>,
<Axes: xlabel='Latitude', ylabel='MedInc'>,
<Axes: xlabel='Longitude', ylabel='MedInc'>,
<Axes: xlabel='MedHouseVal', ylabel='MedInc'>],
[<Axes: xlabel='MedInc', ylabel='HouseAge'>,
<Axes: xlabel='HouseAge', ylabel='HouseAge'>,
<Axes: xlabel='AveRooms', ylabel='HouseAge'>,
<Axes: xlabel='AveBedrms', ylabel='HouseAge'>,
<Axes: xlabel='Population', ylabel='HouseAge'>,
<Axes: xlabel='AveOccup', ylabel='HouseAge'>,
<Axes: xlabel='Latitude', ylabel='HouseAge'>,
<Axes: xlabel='Longitude', ylabel='HouseAge'>,
<Axes: xlabel='MedHouseVal', ylabel='HouseAge'>],
```

```

[<Axes: xlabel='MedInc', ylabel='AveRooms'>,
 <Axes: xlabel='HouseAge', ylabel='AveRooms'>,
 <Axes: xlabel='AveRooms', ylabel='AveRooms'>,
 <Axes: xlabel='AveBedrms', ylabel='AveRooms'>,
 <Axes: xlabel='Population', ylabel='AveRooms'>,
 <Axes: xlabel='AveOccup', ylabel='AveRooms'>,
 <Axes: xlabel='Latitude', ylabel='AveRooms'>,
 <Axes: xlabel='Longitude', ylabel='AveRooms'>,
 <Axes: xlabel='MedHouseVal', ylabel='AveRooms'>],
[<Axes: xlabel='MedInc', ylabel='AveBedrms'>,
 <Axes: xlabel='HouseAge', ylabel='AveBedrms'>,
 <Axes: xlabel='AveRooms', ylabel='AveBedrms'>,
 <Axes: xlabel='AveBedrms', ylabel='AveBedrms'>,
 <Axes: xlabel='Population', ylabel='AveBedrms'>,
 <Axes: xlabel='AveOccup', ylabel='AveBedrms'>,
 <Axes: xlabel='Latitude', ylabel='AveBedrms'>,
 <Axes: xlabel='Longitude', ylabel='AveBedrms'>,
 <Axes: xlabel='MedHouseVal', ylabel='AveBedrms'>],
[<Axes: xlabel='MedInc', ylabel='Population'>,
 <Axes: xlabel='HouseAge', ylabel='Population'>,
 <Axes: xlabel='AveRooms', ylabel='Population'>,
 <Axes: xlabel='AveBedrms', ylabel='Population'>,
 <Axes: xlabel='Population', ylabel='Population'>,
 <Axes: xlabel='AveOccup', ylabel='Population'>,
 <Axes: xlabel='Latitude', ylabel='Population'>,
 <Axes: xlabel='Longitude', ylabel='Population'>,
 <Axes: xlabel='MedHouseVal', ylabel='Population'>],
[<Axes: xlabel='MedInc', ylabel='AveOccup'>,
 <Axes: xlabel='HouseAge', ylabel='AveOccup'>,
 <Axes: xlabel='AveRooms', ylabel='AveOccup'>,
 <Axes: xlabel='AveBedrms', ylabel='AveOccup'>,
 <Axes: xlabel='Population', ylabel='AveOccup'>,
 <Axes: xlabel='AveOccup', ylabel='AveOccup'>,
 <Axes: xlabel='Latitude', ylabel='AveOccup'>,
 <Axes: xlabel='Longitude', ylabel='AveOccup'>,
 <Axes: xlabel='MedHouseVal', ylabel='AveOccup'>],
[<Axes: xlabel='MedInc', ylabel='Latitude'>,
 <Axes: xlabel='HouseAge', ylabel='Latitude'>,
 <Axes: xlabel='AveRooms', ylabel='Latitude'>,
 <Axes: xlabel='AveBedrms', ylabel='Latitude'>,
 <Axes: xlabel='Population', ylabel='Latitude'>,
 <Axes: xlabel='AveOccup', ylabel='Latitude'>,
 <Axes: xlabel='Latitude', ylabel='Latitude'>,
 <Axes: xlabel='Longitude', ylabel='Latitude'>,
 <Axes: xlabel='MedHouseVal', ylabel='Latitude'>],
[<Axes: xlabel='MedInc', ylabel='Longitude'>,
 <Axes: xlabel='HouseAge', ylabel='Longitude'>,

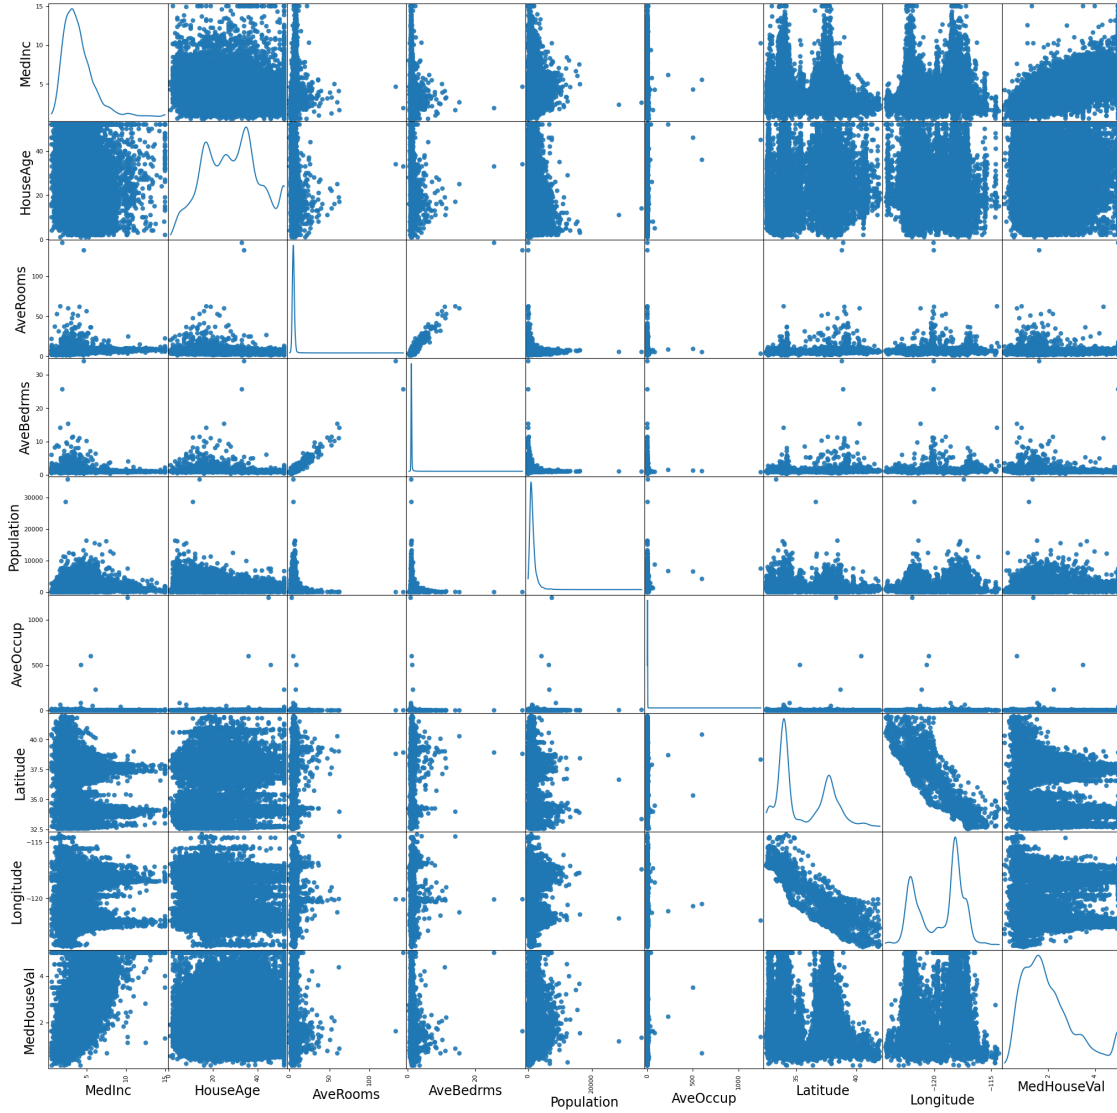
```

```

<Axes: xlabel='AveRooms', ylabel='Longitude'>,
<Axes: xlabel='AveBedrms', ylabel='Longitude'>,
<Axes: xlabel='Population', ylabel='Longitude'>,
<Axes: xlabel='AveOccup', ylabel='Longitude'>,
<Axes: xlabel='Latitude', ylabel='Longitude'>,
<Axes: xlabel='Longitude', ylabel='Longitude'>,
<Axes: xlabel='MedHouseVal', ylabel='Longitude'>],
[<Axes: xlabel='MedInc', ylabel='MedHouseVal'>,
<Axes: xlabel='HouseAge', ylabel='MedHouseVal'>,
<Axes: xlabel='AveRooms', ylabel='MedHouseVal'>,
<Axes: xlabel='AveBedrms', ylabel='MedHouseVal'>,
<Axes: xlabel='Population', ylabel='MedHouseVal'>,
<Axes: xlabel='AveOccup', ylabel='MedHouseVal'>,
<Axes: xlabel='Latitude', ylabel='MedHouseVal'>,
<Axes: xlabel='Longitude', ylabel='MedHouseVal'>,
<Axes: xlabel='MedHouseVal', ylabel='MedHouseVal'>]], dtype=object)

```

<Figure size 1200x1200 with 0 Axes>



1.1.1 1. Residual Sum of Squares (RSS)

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Where:

y_i is the actual value

\hat{y}_i is the predicted value

n is the number of observations

1.1.2 2. Residual Standard Error (RSE)

$$RSE = \sqrt{\frac{RSS}{n - p - 1}}$$

Where:

RSS is the Residual Sum of Squares

n is the number of observations

p is the number of predictors (excluding the intercept)

1.1.3 3. t-statistic

$$t = \frac{\hat{\beta}_j}{SE(\hat{\beta}_j)}$$

Where:

$\hat{\beta}_j$ is the estimated coefficient for predictor j

$SE(\hat{\beta}_j)$ is the standard error of the estimated coefficient for predictor j

1.1.4 4. p-value

$$p = 2 \cdot (1 - T(|t|, df))$$

Where:

t is the t-statistic

df is the degrees of freedom, calculated as $n - p - 1$

T is the CDF of the t-distribution

1.2 Relevant Metrics

Task 1: Fill the missing parts (#TODO) of metric computations (1 Point Each, 3 Points)

```
[10]: def compute_rss(y, y_pred):  
    """  
    Compute Residual Sum of Squares (RSS)  
    y: array of true target values  
    y_pred: array of predicted target values  
    """  
    rss = 0  
    for y_elt, y_pred_elt in zip(y, y_pred):  
        rss += math.pow(y_elt - y_pred_elt, 2)  
    return rss  
  
def compute_rse(y, y_pred, n, p):  
    """  
    Compute Residual Standard Error (RSE)  
    y: array of true target values  
    y_pred: array of predicted target values  
    n: number of observations  
    p: number of predictors  
    """  
    rss = compute_rss(y, y_pred)  
    return math.sqrt(rss/(n-p-1))  
  
def compute_pvalue(X, y, y_pred):  
    """  
    Compute p-values for the coefficients of a linear regression model.  
    X: array of features  
    y: array of true target values  
    y_pred: array of predicted target values  
    return: p-values for each feature  
    """  
    n, p = X.shape # Number of observations (n) and number of predictors (p)  
    # Compute RSS and RSE  
    rss = compute_rss(y, y_pred)  
    rse = compute_rse(y, y_pred, n, p)  
    # Add intercept (constant term) to the design matrix X  
    X = np.c_[np.ones(n), X]  
    # Calculate  $(X^T X)^{-1}$ 
```

```

    XTX_inv = np.linalg.inv(np.dot(X.T, X))

    # Compute standard error (SE) for each coefficient
    se = np.sqrt(np.diagonal(rse ** 2 * XTX_inv))

    # Fit the model to compute the coefficients (betas)
    beta_hat = np.linalg.lstsq(X, y, rcond=None)[0]

    # Compute t-statistics for each coefficient
    t_stats = beta_hat / se

    degrees_of_freedom = n - p - 1

    # Compute p-values
    p_values = 2 * (1 - stats.t.cdf(np.abs(t_stats), df=degrees_of_freedom))

    return p_values

```

1.3 Linear Regression with single predictor

```

[11]: # Select features and target
X = df[['AveRooms']]
#z-normalize the data for each column
X = (X - X.mean()) / X.std()
y = df['MedHouseVal']

# Split the data into training and testing sets (80% training, 20% testing)
# with a fixing seed that ensures same split every time
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

independent_scaler = StandardScaler()
X_train = independent_scaler.fit_transform(X_train)
X_test = independent_scaler.transform(X_test)

# Create a linear regression model
model_1 = LinearRegression()

# Train the model
model_1.fit(X_train, y_train)

# Get the coefficients
print(f"Intercept ( 0): {model_1.intercept_}")
print(f"Coefficients ( 1, 2): {model_1.coef_}")

# Compute RSS for training data
y_pred = model_1.predict(X_train)

```

```

# Compute RSS
rss = compute_rss(y_train, y_pred)

# Calculate R-squared
r_squared_all = model_1.score(X_train, y_train)

# Compute the p-value
p_value = compute_pvalue(X_train, y_train, y_pred)

# Display the coefficients and p-values in a DataFrame
coefficients = np.concatenate([[model_1.intercept_], model_1.coef_])
p_values = np.concatenate([ p_value])

display(pd.DataFrame(pd.DataFrame({'features': ['intercept'] + list(X.columns),
    ↳ 'coefficients': coefficients, 'p-values': p_values})))
print(f"RSS (test data): {rss}")
print(f"R-squared (test data): {r_squared_all}")

```

Intercept (0): 2.071946937378876

Coefficients (1, 2): [0.18323882]

	features	coefficients	p-values
0	intercept	2.071947	0.0
1	AveRooms	0.183239	0.0

RSS (test data): 21518.4672577652

R-squared (test data): 0.025117453148833846

Task 2: Use 'MedInc', 'AveRooms', 'AveOccup', 'HouseAge', 'Latitude', 'Longitude' as predictors.
(2 Points)

```

[12]: X_all = df[['MedInc', 'AveRooms', 'AveOccup', 'HouseAge', 'Latitude',
    ↳ 'Longitude']]
y = df['MedHouseVal']

# Split the data into training and testing sets (80% training, 20% testing)
    ↳ with a fixing seed that ensures same split every time
X_train_all, X_test_all, y_train_all, y_test_all = train_test_split(X_all, y,
    ↳ test_size=0.2, random_state=42)
independent_scaler = StandardScaler()
X_train_all = independent_scaler.fit_transform(X_train_all)
X_test_all = independent_scaler.transform(X_test_all)

# Fit the linear regression model
model_2 = LinearRegression()
model_2.fit(X_train_all, y_train_all)

# Predictions on the test set

```

```

y_pred_all = model_2.predict(X_test_all)

#Code this part
rss = compute_rss(y_test_all, y_pred_all)

# Calculate R-squared
r_squared_all = model_2.score(X_test_all, y_test_all)

# Compute the p-value
p_value = compute_pvalue(X_test_all, y_test_all, y_pred_all)

# Display the coefficients and p-values in a DataFrame
coefficients = np.concatenate([[model_2.intercept_], model_2.coef_])
p_values = np.concatenate([ p_value])

# pd.DataFrame({'features': ['intercept'] + list(X_all.columns), 'coefficients':
↪ coefficients, 'p-values': p_values})
display(pd.DataFrame({'features': ['intercept'] + list(X_all.columns),
↪ 'coefficients': coefficients, 'p-values': p_values}))
print(f"RSS (test data): {rss}")
print(f"R-squared (test data): {r_squared_all}")

```

	features	coefficients	p-values
0	intercept	2.071947	0.000000
1	MedInc	0.708366	0.000000
2	AveRooms	0.045937	0.010799
3	AveOccup	-0.037746	0.000000
4	HouseAge	0.124500	0.000000
5	Latitude	-0.977368	0.000000
6	Longitude	-0.931079	0.000000

RSS (test data): 2259.4509562626686

R-squared (test data): 0.5823077951522642

Task 3: Try model performance on different K values by using the code below, observe the effect of very large K values which one would you pick? (3 Points)

Answer: The optimal value is k=12, since it produces the lowest RSS & highest r-squared output.

```

[13]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

X_all = df[['MedInc', 'AveRooms', 'AveOccup', 'HouseAge', 'Latitude',
↪ 'Longitude']]
X_all = (X_all - X_all.mean()) / X_all.std()
y = df['MedHouseVal']

```

```

# Split the data into training and testing sets (80% training, 20% testing)
↳with a fixing seed that ensures same split every time
X_train_all, X_test_all, y_train_all, y_test_all = train_test_split(X_all, y,
↳test_size=0.2, random_state=42)
independent_scaler = StandardScaler()
X_train_all = independent_scaler.fit_transform(X_train_all)
X_test_all = independent_scaler.transform(X_test_all)

#Fit the KNN model (you can tune 'n_neighbors' for optimal performance)
knn_model = KNeighborsRegressor(n_neighbors=12)
knn_model.fit(X_train_all, y_train_all)

#Make predictions on the test set
y_pred_knn = knn_model.predict(X_test_all)

#Compute RSS and R-squared
rss_knn = compute_rss(y_test, y_pred_knn)
r2_knn = r2_score(y_test_all, y_pred_knn)
print(f"KNN Model RSS: {rss_knn}")
print(f"KNN Model R-squared: {r2_knn}")

```

KNN Model RSS: 1588.7586109921456

KNN Model R-squared: 0.7062949804877341

Task 4: Comment on R-squared and RSS values (1 Point)

R-Squared: This measures which percentage of the total variance in the data is explained by the model, the value lies between 1 & 0, the higher the better. RSS: The residual sum of squares measures how apart the predictions and the true values are, thus a good performing model has a low RSS.

1.4 Visualize results

```

[14]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_squared_error

# Make predictions on the test set for the first model (only AveRooms)
y_pred = model_1.predict(X_test)

# Make predictions on the test set for the second model (multiple features)
y_pred_all = model_2.predict(X_test_all)

# Make predictions using the KNN model
y_pred_knn = knn_model.predict(X_test_all) # Use scaled features for KNN

```

```

plt.figure(figsize=(10, 6))

# Model 1: True vs Fitted (only AveRooms)
sns.scatterplot(x=y_pred, y=y_test, color='blue', label='Model 1 (AveRooms)',
               ↪alpha=0.1)


# Model 2: True vs Fitted (multiple features)
sns.scatterplot(x=y_pred_all, y=y_test_all, color='red', label='Model 2_
               ↪(Multiple features)', alpha=0.1)

# KNN Model: True vs Fitted
sns.scatterplot(x=y_pred_knn, y=y_test_all, color='green', label='KNN Model',
               ↪alpha=0.1)

# Add perfect prediction line
plt.plot([min(y_pred_all), max(y_pred_all)], [min(y_test_all),
               ↪max(y_test_all)], color='black', linestyle='--', label='Perfect Prediction')

# Labels and title
plt.xlabel('Fitted Values (Predicted)')
plt.ylabel('True Values')
plt.title('True vs Fitted Values for Model 1 and Model 2')
plt.legend()
plt.show()

```



Task 5: Compute residuals (1 Point)

```
[15]: ### 2. Residuals vs Fitted

# Compute residuals
residuals_model_1 = y_test - y_pred
residuals_model_2 = y_test_all - y_pred_all
residuals_knn = y_test_all - y_pred_knn

plt.figure(figsize=(10, 6))

# Residuals vs Fitted for Model 1 (only AveRooms)
sns.scatterplot(x=y_pred, y=residuals_model_1, color='blue', label='Model 1_
↳(AveRooms)', alpha=0.1)

# Residuals vs Fitted for Model 2 (multiple features)
sns.scatterplot(x=y_pred_all, y=residuals_model_2, color='red', label='Model 2_
↳(Multiple features)', alpha=0.1)

# Residuals vs Fitted for KNN Model
sns.scatterplot(x=y_pred_knn, y=residuals_knn, color='green', label='KNN_
↳Model', alpha=0.1)

# Add horizontal line at 0 (perfect prediction residual)
plt.axhline(0, color='black', linestyle='--')

# Labels and title
plt.xlabel('Fitted Values (Predicted)')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values for Model 1 and Model 2')
plt.legend()
plt.show()
```